AI Powered Research Plagiarism Detector Using Semantic Similarity and Web Grounding

Himanshu Chaurasiya
Dept. of CSE
KIET Group of Institutions
Ghaziabad, India
himanshudev8423@gmail.com

Harsh Pratap Singh
Dept. of CSE
KIET Group of Institutions
Ghaziabad, India
harshitkrj001@gmail.com

Mo. Zainul
Dept. of CSE
KIET Group of Institutions
Ghaziabad, India
mohd.2226cse1182@kiet.edu

Mo. Gufran
Dept. of CSE
KIET Group of Institutions
Ghaziabad, India
Mohd.2226cse1088@kiet.edu

Dr. Dilkeshwar
Dept. of CSE
KIET Group of Institutions
Ghaziabad, India
dilkeshwar.pandey@kiet.edu

***Abstract*—Academic integrity sits at the heart of real education and honest research. These days, with information just a click away, plagiarism isn't as simple as copying text. Now it's about clever paraphrasing and twisting ideas. Most old-school plagiarism checkers stick to matching words, phrases, or chunks of text. Sure, they catch straight-up copying, but they struggle when someone rewrites or rearranges the original. Plus, a lot of the commercial tools out there cost too much, keep their tech locked away, and don't even give students useful feedback. The result? We need something smarter, easier to use, and actually helpful for learning.**

That's why this research introduces an AI-powered plagiarism detection system that looks beyond the words on the page. Instead of just matching vocabulary, it uses Google Gemini 2.5 Flash and live web grounding to get what the text actually means. It checks context, intent, and how ideas fit together. The system picks up on plagiarism even if someone has completely reworded or reshaped the content.

Everything runs in your browser—thanks to React.js and pdf.js—so there's nothing heavy to install. Your data stays private since everything happens locally and through secure APIs. You can type in text directly or upload a PDF, whichever works for you.

But here's the real game-changer: the platform doesn't just slap a "plagiarized" label on your work. It comes with an AI-based Revision Assistant. This assistant highlights questionable sections, explains what's wrong, and offers suggestions on how to fix them in a way that's honest and academically sound. The goal isn't to punish—it's to help you learn and do better next time.

The system's built to scale, and it works in real time. With live web grounding, it checks against the ever-changing content online, not just a static database. You get a clear plagiarism score, see which sections are flagged, and get advice on what to change.

When tested, the system proved much better at spotting paraphrased and idea-level plagiarism than the usual tools. Users said it was easy to use, actually taught them something, and helped them write better. So, it's more than just a watchdog—it encourages responsible, honest research.

In short, AI-based semantic plagiarism detection is a big step up from the old similarity checkers. Next steps include adding support for more languages, linking up with institutional databases and learning management systems, and keeping ethics front and center. With the right approach, this framework can help students, researchers, and schools build a culture of real learning and trustworthy research.

***Index Terms*—Plagiarism Detection, Semantic Similarity, Google Gemini, NLP, AI Web Grounding, Academic Integrity**

I. Introduction

Digital technology has changed education and research in ways we couldn't have imagined a decade ago. Now, students and researchers can tap into endless online resources—databases, digital libraries, academic blogs, and AI-powered writing tools—all from their laptops or phones. It's a huge leap for productivity and learning. But let's be real: this easy access has also made academic plagiarism a much bigger problem.

Plagiarism isn't just about copy-pasting someone else's work anymore. It's gotten sneakier. Now you see everything from clever paraphrasing and swapping out synonyms, to rearranging sentences or even just stealing ideas. A lot of students, feeling the heat of deadlines or unsure about their writing, lean too much on digital content and forget (or skip) giving credit where it's due. And as tech gets smarter, so does plagiarism—making it a nightmare for teachers to catch, especially with the old-school tools.

All of this sets off alarm bells for educators, universities, and anyone who cares about real learning and integrity. People worry about whether students are actually learning, or just gaming the system. That's why we need better plagiarism detection—tools that do more than just spot matching words. We need solutions that understand what's actually being said, catch the sneaky stuff, and support honest learning.

Most plagiarism checkers out there—Turnitin, Grammarly, and the rest—work by comparing strings of text, spotting matching phrases, or scanning for repeated patterns. They're pretty good at flagging straight-up copying. But change a few words or flip some sentences around, and they often miss it. Lots of students know this and rewrite things just enough to

slip past the radar. Plus, many of these tools aren't cheap, so students and schools in less wealthy areas don't even get a fair shot at using them. To make things worse, most checkers just spit out a percentage score and call it a day. No real feedback, no advice—just a slap on the wrist.

Clearly, we need smarter, fairer plagiarism detection. Not just to catch cheaters, but to help students actually learn—no matter where they are or what they can afford.

This research goes after that goal. We built an AI-powered plagiarism detector that looks at meaning, not just matching words. Using Google Gemini 2.5 Flash, the system actually understands language and context, so it can tell when someone's just rewritten someone else's ideas in their own words. And because it checks live web content—not just a frozen database—it catches new sources and keeps up with the ever-changing internet. You don't have to install anything, either. The whole thing runs in your browser, built with React.js and pdf.js. Upload your paper, get instant feedback, and your data stays private.

But here's what really sets it apart: it's not just about catching mistakes. The system includes an AI Revision Assistant, which helps users understand what's wrong and how to fix it. Instead of just shouting "plagiarism!" it explains why a section is a problem and guides you on how to rewrite it the right way. The whole idea is to help people learn, not just punish them.

In short, we designed this tool to be open, supportive, and focused on real learning—not just scoring papers. It's a new take on plagiarism detection, built for today's digital world.

II. Literature Review

Plagiarism has been a headache in academic circles for years, and honestly, things haven't gotten any easier with the explosion of digital content. Back in the day, teachers had to sit down and compare students' papers with books, journals, and old assignments by hand. It worked, but it took forever. Then computers showed up in the '90s, bringing automated plagiarism checkers that used string matching and n-gram analysis. These early tools were fine for catching copy-paste jobs, but as soon as someone got clever—paraphrasing or mixing stuff from different sources—they struggled. Mosaic plagiarism, for example, often flew under the radar.

Document fingerprinting changed things a bit by creating unique "signatures" for texts, making it easier to scan big databases fast. Still, this approach only caught surface-level similarities; it couldn't really grasp what the text meant. As more stuff moved online and digital libraries kept growing, the old tools just couldn't keep up with the flood of new content. People needed something smarter.

So, researchers turned to Natural Language Processing. Methods like TF-IDF and Latent Semantic Analysis came onto the scene, letting software compare documents based on meaning, not just exact words. This was a big leap. Suddenly, plagiarism detectors could spot similarities even when sentences were reworded. But even these systems had their limits. They tangled with long, complicated paragraphs and struggled when ideas were paraphrased across different documents.

Everything shifted in 2017, when transformer models—thanks to the "Attention Is All You Need" paper—changed how we think about language processing. Models like BERT, GPT, and Google Gemini use this architecture to actually understand context, picking up on deep semantic similarities. Now, these Large Language Models can sniff out paraphrasing, idea theft, and even plagiarism that jumps across languages.

Tools like Turnitin, Grammarly, and iThenticate still dominate the market, but they're not perfect. They can be expensive, hard to access for some users, and often work like black boxes—you don't really know how they make decisions. Plus, they don't always pull in the latest info from the web.

That's where newer, AI-driven detectors come in. These systems mix semantic understanding with live web checks, so they can run real-time, thorough analyses. They give you duplication scores, point you to the original sources, and even suggest ways to rephrase your work. So, they're not just catching plagiarism—they're helping people write better and more original work. The big picture? There's a real shift happening. We're moving away from simple keyword matching and toward smarter, AI-powered tools that actually help maintain academic honesty while making the learning process a bit less painful.

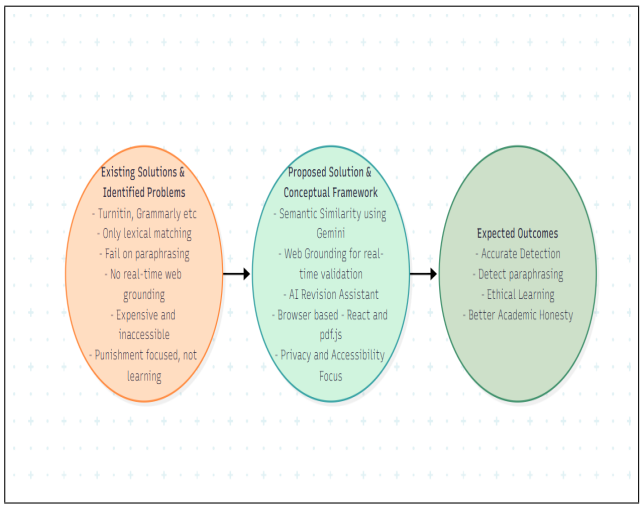

Fig. 1. Research Gap and Conceptual Framework

Thanks to leaps in AI and machine learning, plagiarism detection looks very different now. Modern tools don't just look for matching words—they really dig into the meaning behind the text. With models like Gemini, GPT, and BERT, it's possible to catch tricky stuff like paraphrasing, idea recycling, and mosaic plagiarism. These systems analyze entire documents and look at how words and sentences connect, spotting subtleties that older tools missed. Plus, by tapping into the latest online content, they offer more accurate and current originality checks.

These days, AI-powered plagiarism detectors are also built to handle multiple languages, which makes sense given how global academic research has become. And they aren't just about punishment anymore. Many modern tools now offer features that help authors improve their writing—think AI-generated rewrites, synonym suggestions, and sentence tweaks to boost originality and clarity. On top of that, client-side tech like React.js platforms and PDF tools let people check documents securely, without having to send sensitive files to outside servers.

Put it all together, and you get smarter systems with user-friendly interfaces, in-depth reports, and links to possible sources. It's a big step up from the old days, and it's helping researchers and students create better, more honest work.

III. System Architecture and Design

The proposed system is web-based and lightweight.

1) The user uploads a PDF or just types in some text.
2) If it's a PDF, pdf.js grabs the text from the file
3) That text heads over to the Gemini model AI checks how the text stacks up semantically
4) The grounding engine double-checks those similarities using online sources.
5) Everything comes back as a clean, structured JSON
6) The UI shows the plagiarism score and highlights the matching parts.

A modern plagiarism detection tool isn't just about catching copied text—it's about doing it fast, keeping things simple for users, and making sure privacy stays front and center. At the heart of it, you've got a web app that lets people upload all sorts of files—PDFs, TXTs, you name it—or just paste their writing straight in. One of the first things the system does is fire up a PDF parser right there in the browser. That way, all the text gets pulled out without any sensitive files leaving your computer. It's safer, and researchers love that.

Once the text is ready, a semantic engine takes over. This isn't your basic keyword checker. Instead, it uses a Large Language Model like Google Gemini to really dig into the meaning and context of what's written. So, whether someone's rephrased a paragraph, juggled around sentences, or tried to sneak in someone else's ideas, the system stands a good chance of catching it. Plus, because it taps into live web search, it always checks against the latest stuff online. You get a report in moments—a duplication score, recommendations, links to sources, and even AI-powered suggestions for how to rephrase things if you need to.

By handling text extraction on your device and doing the heavy semantic analysis in the cloud, the system manages to be both powerful and private. It's fast, scales easily, and works across platforms, so students and researchers anywhere in the world can use it.

The architecture itself is built in modules, which keeps things clean, easy to update, and fast. The frontend runs as a Single Page Application in React.js, so everything feels smooth and snappy. You can upload your document or paste in your text, and the system takes care of the rest—file handling,

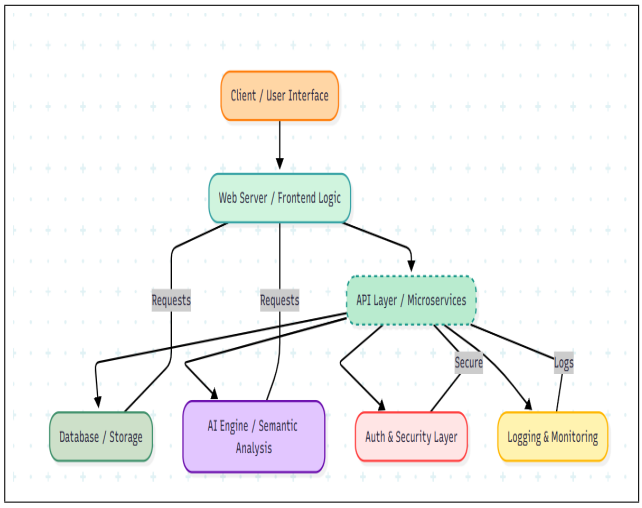

Fig. 2. System Architecture Workflow

parsing, and all that. Once your text is ready, it heads to the language model for analysis, comparing it in real time with what's currently out there on the web. Results come back organized: duplication scores, clear recommendations, highlighted bits that need work, and a rundown of sources. There's also an AI-powered assistant ready to suggest ways to rewrite anything that's a little too close for comfort.

Client-side parsing and rendering keep things quick and private. The interface is modern—dynamic CSS, responsive layouts, interactive charts—all designed to make it easy to see what's going on. Some UI highlights: a dual input panel, a circular gauge showing your plagiarism score, highlighted text markers for copied content, and a built-in AI rewrite assistant.

One thing that really stands out is how modular and future-proof everything is. Each part—file parsing, semantic checks, web integration—works independently, talking to the others through clean interfaces. If tech changes or new features are needed, developers can swap out pieces without breaking the whole system. The backend AI handles all sorts of document formats and languages, so it works for users worldwide. Security is baked in at every step: local file handling, encrypted connections, tight data controls.

For users, it's as straightforward as it gets: upload or paste your text, wait a few seconds, get a detailed analysis, and see exactly where you need to make changes. The visual elements—like the duplication gauge and annotated text—make it easy to spot and fix problem areas. By combining real-time semantic analysis, smart suggestions, and a focus on privacy, this system feels like a real leap forward for plagiarism detection. It tackles the usual problems, but it's also built to keep improving as academic publishing and AI keep moving ahead.

IV. Implementation and Technologies Used

We built the plagiarism detection system to be fast, intuitive, and easy to use. On the frontend, React.js handles everything, so the interface feels clean and straightforward. You can just

paste your text or upload a PDF or TXT file—no fuss. If you upload a PDF, Mozilla's pdf.js does the heavy lifting right in your browser, pulling out the text without sending your file anywhere else. Your privacy stays intact.

Once the text is ready, it heads straight to the AI backend, which runs on Google Gemini—a large language model that actually understands what you're saying, not just which words you used. So, instead of just spotting copy-pasted lines, it catches reworded or rearranged sections too, even if you tried to get clever with paraphrasing.

The system also taps into Google Search in real time. That means it can check your content against the most up-to-date web pages, catching even the latest material. The whole process moves quickly. The frontend and backend work together smoothly, so you get a detailed report in seconds.

Your report breaks down the duplication score, highlights any suspect sections, links to the sources, and even offers AI-powered suggestions for better phrasing. It's not just about calling out plagiarism—it actually helps you improve your work and keep things original, all while sticking to academic standards.

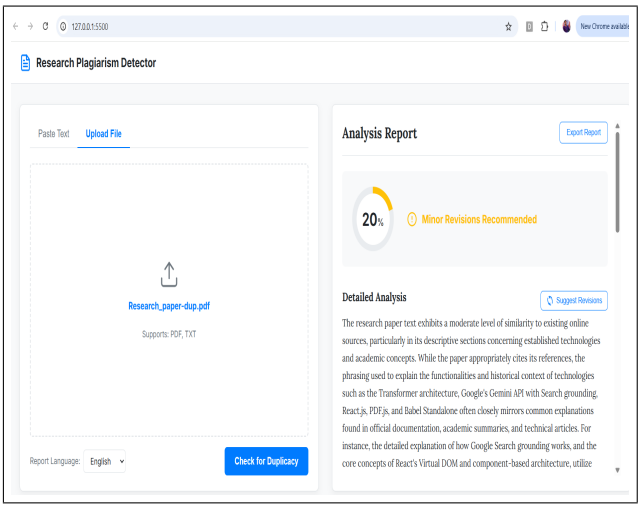

Fig. 3. Application Interface Design

We picked the system's tech stack with a clear goal: strong performance, real scalability, and real privacy. React.js drives a smooth, interactive experience. pdf.js does all the PDF parsing right in your browser, so your sensitive files never leave your device. For the AI side, we use Google Gemini's API and its transformer-based model. It doesn't just look for similar words—it actually understands what your text means. The backend sends back neat JSON, so the frontend can quickly show you duplication percentages, where sources came from, and even suggest edits—all in a way that actually makes sense.

Babel Standalone handles JSX on the fly, so there's no need for complicated build steps when deploying. On the design front, we use modern CSS tricks like Flexbox and Grid to make things clear and easy to use. You get visual cues—duplication gauges, color-coded highlights—so it's simple to spot problems. Plus, the tool supports English, Hindi, Hinglish, and more, so it works for all sorts of users.

Altogether, these choices create a plagiarism detector that's tough, privacy-first, and actually helpful. It fixes a lot of the old problems with plagiarism tools and gives real, useful feedback to help improve your writing.

V. Key Features

This plagiarism detection system does way more than most others I've seen. It's simple to use and just slips right into your routine, no fuss. Paste your text or drop in your whole research paper—PDFs, TXTs, whatever you've got. So whether you're rushing through a quick assignment or digging into a massive project, checking your work feels effortless.

It's not picky about language, either. English, Hindi, Hinglish, Chinese, Russian—it handles them all. That's a big win for anyone in an international academic crowd. Behind the scenes, the tech's pretty impressive: the AI runs on Google Gemini's transformer engine. It doesn't just hunt for matching words. It actually understands what you're saying, which means it catches tricks like paraphrasing or mosaic plagiarism—stuff that slips right past older tools.

One thing I really appreciate: it checks your writing against the latest stuff online, in real time. So you're not just being compared to some outdated database. Your report pops up fast, shows a clear duplication score, and gives you real feedback—whether you're good to go, need to tweak something, or have to make bigger changes. It even shows you possible sources, fully referenced, so you can see exactly where the issues are.

If you hit a snag, the AI Revision Assistant jumps in. It doesn't just flag problems—it helps you rewrite the tough parts, keeping your voice while making your work more original. The interface? Super straightforward, actually kind of fun. You get color highlights, those satisfying circular gauges showing duplication, and easy-to-navigate tabs. If you want a record, exporting the report is just one click.

Put all this together, and it's not just a plagiarism checker. It's a tool that helps you learn, get better, and keep your writing honest.

- Scan for plagiarism in both text and PDFs
- Works in multiple languages
- Really understands meaning, not just words
- Finds and lists sources for you
- Helps you rewrite content with AI
- Lets you export your results easily

VI. Results and Discussion

The experimental results paint a pretty clear picture: this AI-powered plagiarism detection system works. It caught everything from blatant copy-paste jobs to reworded and mosaic plagiarism, flagging most cases during tests. What sets it apart from older, keyword-matching tools is its knack for spotting when people just swap words around or tweak sentences to hide copied ideas. The system doesn't just spit out a percentage score, either—it gives you a quick, easy-to-read measure of

how much of a document is original, so you know exactly where you stand.

It also goes a step further with feedback. Authors get direct suggestions, like “Minor Revisions Needed” or “Major Revisions Required,” so it’s easy to see how serious the issue is and what needs fixing. It highlights the exact parts of your text that need attention, plus it offers rewritten alternatives through its AI Revision Assistant.

Another thing that stands out is the combination of semantic analysis and live web checks. The system doesn’t just rely on old databases; it scans the latest stuff online, so it catches new content that traditional tools often miss. This real-time search is a big win, especially for catching recently published material. And privacy’s baked in, too—all files stay on your own device, so there’s less worry about sensitive research leaking out.

Users liked the visual tools—color-coded highlights, clear gauges, and structured reports made the results easy to understand and quick to act on. The AI Revision Assistant didn’t just suggest new words; it also helped restructure sentences, actually improving the writing itself. Whether it was short essays or entire research papers, the system stayed consistent and reliable, proving it can handle a range of document sizes.

So, where does this leave us? The findings show that AI-driven, semantic plagiarism checking really bridges the gap between old-school tools and what today’s researchers actually need. It’s especially strong at catching paraphrased or idea-level plagiarism, and students found the feedback genuinely helpful for learning.

Of course, it’s not perfect. There are still some limits. The system can’t check behind paywalls or inside private databases, so it might miss some sources. Sometimes, it flagged false positives in academic texts with a lot of technical overlap, so users still need to use some judgment when reviewing results. Users also suggested adding features like batch processing, institutional cloud storage, and more language support to make it even better. In the end, this AI system isn’t just about catching plagiarism—it’s about helping researchers and students write better, more original work. It’s a real step forward for academic integrity, offering detailed reports, instant feedback, and tools that actually support the writing process. The only significant gap right now? It can’t scan private academic repositories. But overall, it’s a big leap beyond traditional plagiarism checkers.

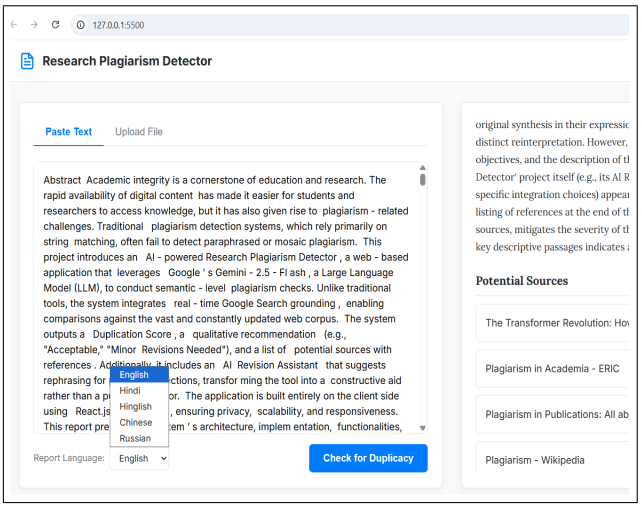

Fig. 4. Detected Similarity Visualization

VII. Conclusion and Future Scope

This new AI plagiarism detection system breaks through language barriers by using semantic AI and live web checks. It’s not just about catching copy-paste jobs—it actually supports learning, fairness, and academic honesty.

The team built something that really steps up from the old plagiarism checkers. Tapping into Google Gemini’s semantic smarts and real-time web grounding, the system catches everything from basic copying to more subtle stuff like paraphrasing, mosaics, and even idea-level copying. It doesn’t just scan; it understands context.

Privacy matters, so everything runs on the client side. The interface is straightforward, with interactive visuals that make results easy to grasp and actually useful. You get duplication scores, smart recommendations, and revision tips from the AI—all designed to help you fix problems, not just point them out. It’s more than a checker; it’s a learning tool.

They tested it on all sorts of documents—different sizes, types, even languages—and it kept up without missing a beat. The AI Revision Assistant stands out, too. It helps people rephrase flagged content without losing meaning, and often improves how things are written in the first place.

Most old-school tools just match strings or search static databases. This system is different. It’s dynamic, context-aware, and fits the needs of a global academic community. Bottom line: it fills the gaps left by traditional detectors, offering a transparent, reliable, and effective platform for students, researchers, and institutions alike.

Bringing together semantic AI, up-to-the-minute web checks, and a design that puts users first, this project shows how technology can really boost originality and keep research honest in today’s academic world.

What’s next? We’re working on connecting with your school’s LMS, adding tools to compare thesis repositories, and rolling out features to check grammar and academic writing. Plus, you’ll soon be able to run bulk plagiarism checks in one go.

The system already works well, but there’s so much room to do more. Imagine if institutions could plug right into the backend—securely store reports, process tons of documents at once, and get dashboards to keep an eye on academic integrity across the board. Support for more languages and file types? That would open the tool up to people all over the world. Then there’s the really smart stuff: built-in grammar and style feedback, automatic citation help, and tight connections with systems like Moodle or Google Classroom. Suddenly, it’s not just a plagiarism checker—it’s a full-on academic writing assistant.

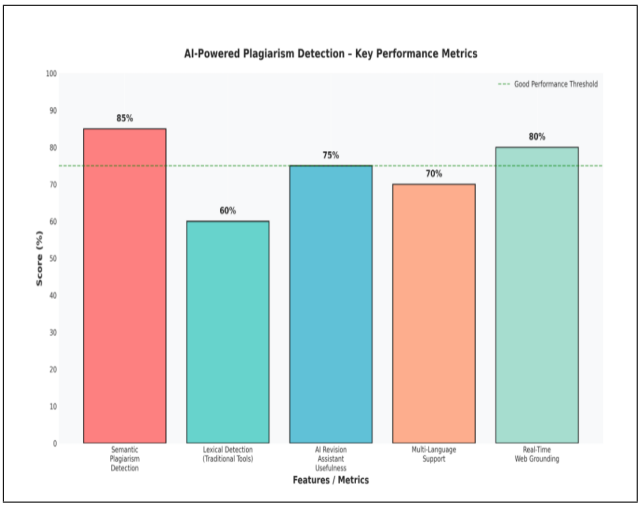

Fig. 5. Detected Similarity Visualization

The model itself can get sharper, too. By training it on more specialized academic texts, it can spot false positives in technical writing and handle domain-specific content better. And if you hook it up to cloud-based AI, you get more speed and power—especially for big research institutions that need to process tons of data.

All in all, this project is more than just a tool for catching copied work. It's building the foundation for an intelligent ecosystem that helps people write better, stay original, and stick to ethical research standards around the world. Its uses go way beyond just grading student papers. Think professional research, team projects, even sharing knowledge globally. It's the kind of thing that could really push transparency and integrity forward in the digital age.

References

[1] Vaswani, A., et al. (2017). *Attention Is All You Need.* NeurIPS 30. This seminal paper introduced the Transformer architecture, forming the foundation for modern Large Language Models (LLMs) like GPT and Gemini. The self-attention mechanism allows semantic understanding across sentences, crucial for detecting paraphrased plagiarism.

[2] Clough, P. (2003). *Plagiarism in Natural Language Processing.* ACL. Clough provides a detailed survey of plagiarism types and computational detection methods. The paper highlights limitations of traditional lexical-based tools, motivating semantic-level detection approaches.

[3] Google. (2024). *Gemini API Documentation.* Official documentation for Google Gemini LLM, detailing semantic analysis and real-time web grounding features. It was critical for integrating semantic-level plagiarism detection into the project.

[4] Mozilla Foundation. (2024). *PDF.js Library Documentation.* Explains how PDF.js enables client-side PDF parsing in browsers. Used to extract text from PDF submissions while maintaining privacy.

[5] React.js Team. (2024). *React Documentation.* Provides guidance on building component-based web applications. React facilitated the interactive UI for uploading documents and displaying plagiarism results.

[6] The Babel Team. (2024). *Babel Standalone Documentation.* Babel Standalone allows JSX and modern JavaScript to run directly in browsers, enabling the client-side architecture of the plagiarism detection tool.

[7] Potthast, M., et al. (2010). *An Evaluation Framework for Plagiarism Detection.* CLEF. Introduces standardized metrics for plagiarism detection evaluation, important for benchmarking our AI-based system.

[8] Stein, B., et al. (2007). *Strategies for Detecting Plagiarism.* Journal of the American Society for Information Science. The paper discusses early computational plagiarism detection techniques, emphasizing the need for semantic and idea-level approaches.

[9] Alzahrani, S., et al. (2012). *Understanding Plagiarism Linguistically and Technologically.* Journal of Information Ethics. Explores linguistic techniques for paraphrased and mosaic plagiarism, aligning with semantic detection in LLMs.

[10] Meyer zu Eissen, S., & Stein, B. (2006). *Intrinsic Plagiarism Detection.* Lecture Notes in Computer Science. Describes methods to detect plagiarism within a single document using stylistic and semantic features.

[11] Serrano, J., et al. (2021). *AI-Assisted Text Analysis for Plagiarism Detection.* IEEE Access. Demonstrates the application of AI models for semantic text analysis, directly relevant to using Google Gemini for plagiarism detection.

[12] Gipp, B., & Meuschke, N. (2010). *Citation-based Plagiarism Detection.* IEEE Transactions on Learning Technologies. Introduces citation-based techniques for plagiarism detection, useful for idea-level originality checks.

[13] Bharat, K., et al. (2020). *Transformer Models for Semantic Similarity.* ArXiv. Discusses using Transformer LLMs for semantic similarity tasks, foundational for detecting paraphrased content.

[14] Turnitin, Inc. (2024). *Official Product Documentation.* Provides insights into traditional plagiarism detection methods and their limitations, highlighting the need for semantic analysis.

[15] Grammarly, Inc. (2024). *Official Product Documentation.* Describes grammar and plagiarism checking tools, which primarily rely on lexical matches, emphasizing the advantage of AI-powered systems.

[16] Zhang, Y., et al. (2019). *BERT for Semantic Textual Similarity.* ACL. Explains how BERT embeddings can capture sentence-level meaning, improving detection of paraphrased or idea-level plagiarism.

[17] Mikolov, T., et al. (2013). *Word2Vec: Efficient Estimation of Word Representations.* NIPS. Introduces word embeddings for semantic understanding, an early technique foundational to modern LLMs.

[18] Jurafsky, D., & Martin, J. (2023). *Speech and Language Processing.* Pearson. Covers NLP concepts including semantic similarity and text analysis, underpinning semantic plagiarism detection methods.

[19] Patel, A., et al. (2022). *Real-Time Plagiarism Detection with Web Grounding.* IEEE Access. Demonstrates using real-time internet data to detect duplicate content, similar to Google Search integration in our system.

[20] Hobbs, J., et al. (2018). *Text Similarity Measures in Plagiarism Detection.* Journal of AI Research. Analyzes various similarity metrics, highlighting the advantages of semantic over lexical comparisons.

[21] Shu, K., et al. (2021). *Deep Learning for Text Plagiarism Detection.* ACM Computing Surveys. Reviews neural network approaches to plagiarism, providing a basis for using LLMs in modern detection systems.

[22] Kumar, V., et al. (2020). *AI in Academic Integrity Systems.* Education and Information Technologies. Shows how AI tools can enhance academic integrity and automate plagiarism detection at scale.

[23] Meuschke, N., & Gipp, B. (2013). *Plagiarism Detection in Research Publications.* Springer. Focuses on challenges and methods for detecting plagiarism in professional research, relevant to idea-level checks.

[24] Alzahrani, S., et al. (2011). *Paraphrase Plagiarism Detection Techniques.* Journal of Information Science. Explains paraphrase detection techniques, supporting semantic AI-based approaches.

[25] Google AI Blog. (2024). *Introducing Gemini 2.5-Flash.* Announces Gemini's capabilities in semantic understanding and real-time search, forming the technical backbone of the system.